

Search coverage and remaining work

September 7-8, 2026 | Historical discovery campaign and documented coverage limits

Historical discovery report. Current Core, Archive and Borderline decisions are recorded in the [September 8 relevance audit](#). The totals and stopping-rule result below use the broader inclusion boundary in force during that campaign.

Campaign: 7-8 September 2026. The revised search stopping rule was **met on 8 September 2026**: rounds 17 and 18 each completed the fixed queries and meaningful source-link reviews, with no new qualifying project or component after canonical reconciliation. The final baseline contains 6,521 units. The validated registry update contains 6,227 active cards, including 6,051 new cards; its workbook has 6,228 rows including the historical held record. Publication artifacts passed the final record, reference, workbook and application checks; deployment status is tracked separately. This is a bounded process result, not exhaustive coverage of the public ecosystem.

The earlier published registry contained 177 cards. Its selected update added 45 entries, held 7 and excluded 7. That was a reviewed batch, not completion of the full discovery backlog. Availability checks of published URLs were separate from capability review.

The continuation completed a revised, source-specific search plan. It did **not** complete the original broad matrix of every planned technology term, workflow term and source. Amendments expanded ambiguous acronyms and narrowed unfinished homonyms. Amendment 4 then changed remaining repository searches from name/description/README matching to name/description/topics after a precision check showed substantial unrelated results. Completed broad searches and partial observations were retained. This improves the relevance of the revised search but reduces README-only discovery coverage; it is a post-start scope change, not retrospective completion of the original plan.

The fixed repository partition corpus has a cutoff of **7 September 2026**. Fresh queries run on 8 September repeat that bounded corpus; they do not cover repositories created after the cutoff. Repeated query dates therefore do not extend the repository-partition lane's creation-date boundary. Smithery, code and indexed-directory lanes retain their own recorded search rules and retrieval dates. The retained Official Registry harvest log is dated 7 September at 06:21 UTC; later phrase screens reuse those bytes, rather than retrieving a new live inventory. The retained Docker checkout is commit 8c773729f13f036da8c909be503fe433923a9aa2, dated 21 August 2026. That commit date describes the retained source version, not a fresh harvest date. Individual retrieval timestamps remain preserved.

Source before expansion rounds	Recorded coverage	What the count means
GitHub repositories	277 revised queries; 700 result pages; 51,103 observations	152 completed and 125 inherited completed query records; no recorded page-sequence gaps or incomplete-result flags in this revised set. Observations overlap.

Source before expansion rounds	Recorded coverage	What the count means
GitHub skill/code search	16 queries; 53,965 observations	Collector-reported completion. Sharded logs preserve size partitions and results but lack per-page records, so page-by-page completeness cannot be independently verified. Indexed code is limited to 393,216 bytes and is not all repository content.
Smithery	214 queries; 389 pages; 4,925 unique listings	Vocabulary searches with pagination, not a complete unfiltered directory census.
Indexed Glama, PulseMCP and MCP.so results	48 queries: 16 per domain	Search-engine results, not full directory enumeration or exhausted result pagination.
Official MCP Registry	28,167 latest-version records across 282 reported pages	Live pagination rather than a transactionally frozen snapshot. Automated metadata screening does not mean every record received individual primary review.
Docker registry checkout	496 files	A retained repository snapshot screened locally; 496 is a file count, not a server count.

Round-specific counts and stopping statements below describe the evidence available when each round closed; the final result is stated above.

Historical attempts remain visible: the query-status ledger contains 19 errors, 2 interruptions, 9 incomplete attempts and 1 interruption for a prospective scope amendment, alongside completed and inherited records. Revised successful queries do not erase these failures or complete the original expressions. The accompanying audit summary preserves each unsuccessful/interrupted query and its status. Earlier capped searches and superseded batches are historical evidence, not additional completed searches or new candidate counts.

Expansion round 1

Expansion round 1 completed 60 repository searches, 20 code searches, 20 Smithery searches and 60 indexed searches, followed by source and link review. Its code lane retrieved 1,627 observations and reviewed 229 distinct new source bodies. Across the page-link closure,3,073 observations received a disposition. This was a **positive-yield round**. The canonical comparison recorded 117 newly represented eligible units and 107 new project identities; aliases and component regrouping affect these figures, so they are not counts of new public cards.

The PopHIVE MCP wrapper was recovered in a recorded four-query expansion batch. This establishes wrapper recovery at batch level, not attribution to an individual query or a formal pass of the original PopHIVE-project control. The four exact queries and domain restrictions are retained in the accompanying audit summary.

Expansion round 2

Round 2 also closed with qualifying discoveries. It completed 60 repository queries, 20 code queries, 20 Smithery queries and 60 indexed queries. Indexed searching produced 75 focal pages and 455 directory identity observations; 71 promoted repository candidates received a primary-source disposition. The code lane retrieved 2,924 observations and reviewed 646 distinct new source bodies: 42 eligible, 313 discovered components not separately approved, 278 excluded, 9 held and 4 duplicates. These are source-body dispositions, not public cards. All 42 eligible code bodies received outgoing-link closure.

The round 2 canonical comparison added 112 units and 109 project identities, with no unresolved conflicts. A round 3 baseline was then frozen at 5,867 units. Neither of those two expansion rounds was a zero-yield round, so the consecutive-zero counter remains zero. These canonical figures precede final publication grouping and do not replace the 177-card historical baseline.

Round 2 also screened 20 alternative phrases against retained Official Registry and Docker snapshots, finding no matches. These were automated phrase screens, not newly harvested inventories. The Docker screen records 569 tracked files; the initial methods recorded 496 checkout files. Those different recorded file denominators are retained separately and must not be added or presented as server counts.

Expansion round 3

Round 3 is also closed and positive. It completed 60 repository queries, 20 code queries, 20 Smithery queries and 60 indexed-directory queries. The indexed lane retained 103 focal pages: 97 returned HTTP 200, four returned 404 and two returned 403. Failed pages remain explicitly unavailable. Its 461 linked identity observations comprise 142 prior matches, 231 scope/context negatives, 67 primary-review candidates and 21 navigation links. The 67 primary candidates received 31 eligible, 26 excluded, seven held, two duplicate and one discovery-collection disposition. Two hosted primary components also qualified: Hex's analytical workspace interface and a recovered first-party Olyport MMWR bibliographic tool contract.

The repository lane dispositioned 317 identities, including 153 explicit human-skill homonyms and other automated or contextual negatives; these are not 317 full-document reviews. The code lane reviewed or reconciled 254 distinct new bodies from 429 observations: 32 eligible, 119 discovered but not separately approved, 76 excluded, 24 held, two duplicates and one discovery collection. All 32 eligible bodies' outgoing links were closed. A separate directory follow-up covered 52 entries and 53 linked archived variants, with no approvals. Eleven explicitly attributed copies also received original-source reconciliation: nine narrowly approved original components, one duplicate and one held unavailable original; the nine approvals belong to existing repositories and do not approve the copies or whole collections. The Smithery lane found 1,912 identities, including 48 new ones; 47 were metadata-scope negatives and the remaining Hex listing was reconciled to its primary review.

After source-link and identity reconciliation, round 3 added 135 canonical units and 121 project identities. The round 4 baseline is frozen at 6,002 units. These are provisional editorial units, not public-card counts. All three completed rounds have positive yield, so the consecutive-zero counter remains zero. Link-role counts include repeated documentation, upstream APIs and assets: 702 of the 1,334 non-GitHub URL observations were repository tree/blob provenance URLs rather than new project pages.

Expansion round 4

Round 4 is closed and positive. It completed 60 repository queries over 86 pages, 20 code queries, 20 Smithery queries over 40 pages, and 60 indexed-directory queries. The indexed lane retained 55 focal pages: 54 returned HTTP 200 and one returned 404. Its 61 primary projects received 43 eligible, 11 excluded, six held and one duplicate disposition. The unavailable page remains a recorded coverage limitation.

The repository lane returned 7,058 observations and 6,239 distinct identities. These received 126 eligible, 6,076 excluded, 18 held, 17 duplicate and two discovery-collection dispositions. The review denominator is 240 individual dispositions, 5,993 automated metadata negatives and six automated canonical matches. Search stemming around "QALY" produced many unrelated "QA" results; these explicit automated negatives are not described as manual document reviews. The Smithery lane returned 1,902 identities, with 47 previously unseen listings: one narrowly eligible air-quality tool component, one held statistics service and 45 exclusions. The code lane covered 115 distinct new bodies from 177 observations: 26 eligible, 59 discovered but not separately approved, 16 excluded, 12 held and two duplicates.

All promoted primary and scoped source-link followups are closed. Linked documentation and assets do not become approved project cards merely because they were encountered. Round 4 added 227 canonical keys and 206 project identities, with two removed or regrouped keys, for a net increase of 225 units. The round 5 baseline is frozen at 6,227 provisional editorial units. All four completed expansion rounds have positive yield, so the consecutive-zero counter remains zero. Fresh fixed-union search rounds begin from this baseline; earlier exploratory rounds remain separately documented.

The prior published site registry and published downloadable workbook each contained 177 historical entries. Git-history inspection confirmed the published workbook at both the GitHub and deployed-site revisions. An older local 28-entry template was used during reconstruction; it was not the previously published workbook. New public cards and the rebuilt workbook require their own final counts and checks.

Expansion round 5

Round 5 completed the fresh fixed-union search plan and remained positive after a pagination audit. It ran 240 repository queries, 80 code queries, 80 Smithery queries and 240 indexed-directory queries. Repository searches returned 6,602 identities: 6,597 prior matches and five new identities resolved as four exclusions and one exact-body/canonical duplicate. Smithery returned 3,647 identities; its three new listings received two exclusions and one hold. An additional empty Smithery boundary page increased the retained page count to 160.

The code-search audit found that reported result totals could underestimate the retrievable result set. Full terminal pages therefore required continuation even when the reported total had been reached. Those continuations yielded additional source bodies and new qualifying research components. Repository terminal pages were also audited, with no missing full-page boundaries found. These corrections are retained in the audit history; the provisional zero-yield result was not accepted.

The indexed lane retained 250 focal pages: 240 returned HTTP 200, five returned 404 and five returned 403. Its 1,383 linked identities comprised 543 prior canonical matches, 821 prior directory scope dispositions and 19 novel identities. Four new primary documents were excluded, 14 were scope or supporting-reference negatives, and one encoded identity matched a prior project. All source and link followups are closed. After final reconciliation,

round 5 added 7 canonical keys and 4 project identities; the next baseline contains 6,234 units. The consecutive-zero counter remains zero, and two sequential zero-yield rounds are still required.

Expansion round 6

Round 6 is also closed and positive under the fixed-union plan: 240 repository queries, 80 Smithery queries, 240 indexed-directory queries and 80 code queries. The indexed lane retained 249 focal pages, with 239 HTTP 200 responses, five 404 responses and five 403 responses. Its 1,391 linked identity observations included 533 prior canonical matches, 829 prior directory dispositions and 29 novel identities. Ten promoted primary reviews yielded two narrowly eligible components, four holds and four exclusions; 19 other identities received contextual negatives.

The two primary approvals cover medical-document text/list retrieval and named-anatomical image-segmentation orchestration. Their limitations explicitly exclude unsupported vector-ranking claims and clinical accuracy guarantees. A separate hosted approval covers a small, version-linked research-paper passage collection; it does not approve benchmark task execution or scoring. All source-link followups are closed. Final reconciliation added 4 canonical keys and 4 project identities, bringing the next baseline to 6,238 units. The consecutive-zero counter remains zero.

Expansion round 7

Round 7 is closed and positive. It repeated the fixed 240 repository, 80 Smithery, 240 indexed-directory and 80 code queries. The indexed lane retained 253 focal pages: 243 returned HTTP 200, five returned 404 and five returned 403. Eight initial timeouts were successfully retried, with both attempts retained. Among 1,401 linked identity observations, 542 matched prior canonical identities, 846 matched earlier directory dispositions, and 13 were novel. Five primary reviews produced four eligible statistical-retrieval components and one exclusion; eight other identities were contextual negatives.

The four approvals cover German and North Macedonian official-statistics packs within the shared Pipeworx gateway, a Finland PxWeb adapter and the core Eurostat retrieval branch of a larger suite. These are bounded adapter approvals, not approval of all gateway routes or every suite module. A bounded project-link follow-up examined the Pipeworx public organization catalogue: 1,345 repositories across 14 pages. It reconciled 76 prior adapters and two already reviewed packs, screened 1,149 repositories outside this particular family scope, and individually reviewed 118 additional official-statistics/open-data packs. Those reviews produced 95 narrowly eligible packs, 22 exclusions and one hold. All 725 outgoing source-link observations were closed. These packs share infrastructure and are not 95 independent statistical methods. The code and all source-link followups are closed. Final reconciliation added 100 canonical keys and 100 project identities, bringing the next baseline to 6,338 units. The consecutive-zero counter remains zero.

Expansion round 8

Round 8 is closed and positive. It repeated 240 repository, 80 Smithery, 240 indexed-directory and 80 code queries. The indexed lane retained 252 focal pages: 242 returned HTTP 200, five returned 404 and five returned 403. Its 1,382 linked identities included 543 prior canonical matches, 816 prior directory dispositions and 23 novel identities. Two initial primary reviews produced one narrowly eligible SymMap association-retrieval component and one administrative fraud-scoring exclusion; 21 other identities received contextual scope exclusions.

New academic tag and recommendation links required 21 additional primary reviews, yielding 14 eligible observations, three holds, two exclusions and two duplicate listings. These reviews include bounded scholarly retrieval, document preparation, research-figure drafting and sensitive-data redaction workflows. The SymMap approval excludes automatic evidence-tier validity and any treatment-effect inference; image-generation approvals do not validate numerical plots or journal compliance. Source-link observations were reconciled separately from primary reviews, including 34,489 non-GitHub observations from the focal pages.

Smithery's six new listings produced five exclusions and one newly documented method-catalogue filtering component within the existing Agentic RL project. The Alpha Sophia query playbook likewise belongs to an existing hosted project. The code lane completed 80 queries over 216 pages, returning 8,403 observations: 8,384 exact prior observations and 19 newly reviewed bodies. These yielded no eligible bodies: ten exclusions, one hold, three discovered-only variants and five duplicates. After all primary and scoped link followups closed, canonical reconciliation added 16 units and 14 project identities. The next immutable baseline contains 6,354 units. The consecutive-zero counter remains zero; no stopping-rule claim is made.

Expansion round 9

Round 9 is closed and positive. It repeated 240 repository, 80 Smithery, 240 indexed-directory and 80 code queries. The indexed lane retained 246 focal pages: 236 returned HTTP 200, five returned 404 and five returned 403. Its 1,372 linked identities included 544 prior canonical matches, 813 prior directory dispositions and 15 novel identities. Two initial primary reviews produced one narrowly eligible USGS watershed-attribute lookup and one historical SimPy package reconciliation; the other 13 novel identities received contextual scope exclusions.

Following the new project links required eight additional primary reviews: four eligible and four excluded. The approved scopes are documented queue simulation and analytic comparison, a limited emergency-department demo, Weibull distribution summaries, and AnyLogic model preparation and structural checking. These approvals do not establish clinical validity, general model accuracy or access to undocumented tools. The industrial teaching demos and generic financial or business calculators were excluded. Source-link roles include 34,027 non-GitHub observations from the focal pages; repeated URLs and navigation were not counted as individual primary reviews.

The repository and Smithery lanes found no new eligible sources. The code lane completed all 80 queries over 216 pages, returning 8,405 observations: 8,403 exact prior observations and two newly reviewed bodies. Both bodies were discovered-only variants, with no new eligible code approval. Separate catalogue references to unavailable ITIS and ORCID components remained held, and all source-link followups closed. After all lanes closed, canonical reconciliation added 5 units and 5 project identities. The next immutable baseline contains 6,359 units. The consecutive-zero counter remains zero; two subsequent completed zero-yield rounds are still required.

Expansion round 10

Round 10 is closed and positive. It repeated 240 repository, 80 Smithery, 240 indexed-directory and 80 code queries. The indexed lane retained 242 focal pages: 232 returned HTTP 200, five returned 404 and five returned 403. Its 1,349 linked identities included 531 prior canonical matches, 795 prior directory dispositions and 23 novel identities. Four initial repository reviews produced one narrowly eligible retrieval component and three exclusions; another apparent repository was correctly identified as a Wikipedia integration directory. The remaining 18 novel identities received contextual scope exclusions.

The resulting source chains required 63 editorial records, with distinct denominators: four initial repository reviews and one integration-directory identity review, 19 additional project reviews, three directory-purpose screens, one unavailable collection metadata disposition, three same-project skill/context reviews, and 32 component dispositions within the previously reviewed Stall project. These records yielded 32 eligible decisions, 19 exclusions, nine holds, two duplicates and one discovery collection. The Stall review approved only 20 specific components, held seven and excluded five; a listing of 210 tools did not receive blanket approval. Polyform likewise counts as one hosted primary project, with 67 contract screens and 21 narrowly approved data subsets.

New approvals include constrained environmental and inflation context retrieval, regulatory-source discovery, scholarly retrieval, document preparation and source-annotation access. Limitations exclude inferred treatment effects, unvalidated hazard or compatibility judgments, undeclared provider coverage and inaccessible or retired interfaces. One community registry failed two metadata retrieval attempts and remains held unavailable, rather than being described as inspected. The 33,574 non-GitHub observations from focal pages were reconciled separately from primary reviews.

The repository and Smithery lanes found no new eligible sources. The code lane completed 80 queries over 218 actual pages, returning 8,404 observations: 8,403 exact prior observations and one newly reviewed evidence-synthesis instruction body, which qualified narrowly. It had no outgoing source links. After all lanes and scoped project links closed, canonical reconciliation added 33 units and 13 project identities. The next immutable baseline contains 6,392 units. The consecutive-zero counter remains zero; two subsequent completed zero-yield rounds are still required.

Expansion round 11

Round 11 is closed and positive. It repeated the same 240 repository, 80 Smithery, 240 indexed-directory and 80 code queries. Repository search returned 7,442 observations across 318 pages; Smithery returned 13,454 across 160 pages. Neither lane added an eligible source. The code lane inspected 218 actual result pages and reconciled all 8,403 observations to prior source bodies, with no novel body. Indexed discovery retained 250 focal pages: 240 returned HTTP 200, five returned 404 and five returned 403. Its 1,369 linked identities included 523 canonical matches, 777 previous directory dispositions and 69 novel identities. Twelve initial editorial records produced six eligible decisions, while 57 other novel identities received directory-context scope exclusions.

Source chains from those discoveries required separate, bounded collection work. Minia's 1,678 metadata entries were screened, and 188 distinct selected method or service contracts were reviewed individually: two eligible, 132 held for specific contract gaps, 53 excluded and one duplicate. These holds describe missing method, input or result documentation; they are not 132 failed runtime tests or separate public project cards. StackResolve's 416 profiles received 383 contextual screens and 33 primary reviews, with nine additional linked component reviews. The primary reviews yielded 19 eligible, seven excluded, two held and five duplicates; the linked reviews yielded four eligible, two excluded, one held and two duplicates.

A separate Zinin review covered 25 bounded component records, with 21 eligible and four held, including explicit conflicts between current actor documentation and older gateway claims. Thirty-nine other linked Apify records yielded 21 eligible, 17 exclusions and one duplicate. Pure debt, business confidence, construction, freight and industrial datasets were excluded despite official-source branding. Additional automatic platform recommendation widgets were recorded as navigation; they did not trigger an unrestricted platform census or receive approval.

Authored links and selected source contracts were followed separately. The 35,799 non-GitHub focal observations were reconciled as URL roles, not counted as individual primary reviews.

After source and link closure, round 11 added 75 canonical units and 66 project identities, producing an immutable next baseline of 6,467 units with no removed units or unresolved adverse conflicts. All 177 historical identities remain represented. The consecutive-zero counter remains zero. Round 12 begins from this frozen baseline; the required two consecutive completed zero-yield rounds have not occurred.

Expansion round 12

Round 12 is closed and positive. The indexed lane completed 240 queries and retained 248 focal pages: 238 returned HTTP 200, five returned 404 and five returned 403. All 1,380 linked identities matched previous canonical or directory dispositions. The 34,348 non-GitHub observations were reconciled, including ten new URL contexts that supplied no new qualifying primary lead. Five matches from the retained Official snapshot matched prior bounded identities; the Docker phrase screen found none. Repository and Smithery searches also closed without new eligible sources: 240 repository queries returned 7,439 observations across 318 pages, and 80 Smithery queries returned 13,454 observations across 160 pages.

Late in the code lane, two previously unseen source bodies appeared. All 80 code queries completed over 217 actual pages, returning 8,403 observations: 8,401 exact prior observations and two new bodies. A Vietnamese literature-planning instruction component qualified narrowly; a generic competitive-programming planning component was excluded. Six pinned, source-linked guides were then reviewed separately: five qualified for bounded research preparation and one was held. No external links remained unresolved. These are documented planning components, not claims that database searches or validated statistical analyses were executed.

The final canonical comparison adds six components in one new project, increasing the immutable baseline from 6,467 to 6,473 units with no removals or unresolved conflicts. The consecutive-zero counter remains zero. Zero results in the other lanes do not make the whole round a zero. Round 13 begins only after this closed positive round and verified frozen baseline; two consecutive completed zero-yield rounds remain required.

Expansion round 13

Round 13 is closed and positive. All 240 repository, 80 Smithery, 240 indexed-directory and 80 code queries completed. The indexed lane retained 249 focal pages: 239 returned HTTP 200, five returned 404 and five returned 403. Of 1,360 linked identities, 528 matched canonical records and 820 matched previous directory dispositions; 11 new identities received secondary purpose exclusions and one sports-odds adapter received a primary-source exclusion. Its shared gateway advertisements did not extend approval to unrelated methods.

The 34,534 non-GitHub observations led to one new bounded hosted research enabler: Kinetic's preference-study preparation and access to account-owned respondent data and stored results. Approval does not validate its proprietary estimation, confidence labels, health-state utilities or commercial revenue forecasts, and does not cover its full 100-tool surface. All 298 reviewed source-link contexts were closed. A separate Patsnap package was reconciled as a duplicate: different endpoint and tool names did not establish new research scope beyond existing patent retrieval.

The code lane completed 213 actual pages with 8,046 observations, including one new generic engineering body that was excluded. Repository and Smithery lanes added no eligible sources. Canonical reconciliation therefore

adds exactly one component/project, bringing the next frozen baseline to 6,474 units without removals or unresolved conflicts. The consecutive-zero counter remains zero; the stopping rule is still unmet.

Expansion round 14

Round 14 is closed and positive. It completed the same 240 repository, 80 Smithery, 240 indexed-directory and 80 code queries. Indexed discovery retained 253 focal pages: 243 returned HTTP 200, five returned 404 and five returned 403. Among 1,382 linked identities, 529 matched canonical records and 841 matched previous directory dispositions. Four new primary sources qualified narrowly; eight other identities received contextual exclusions. A further hosted source qualified, another was excluded, and the qualifying source's skill/stdio wrapper was reconciled as duplicate packaging.

The five new approvals cover selected CMS and NHS data access, identifier normalization/checksums, supplied-data scientific chart preparation, and observed job-posting/salary/count retrieval. Their boundaries exclude patient-level or clinical validity, proof that an identifier was assigned to a person, unvalidated numerical conclusions, and treating a curated posting sample as population employment or hiring incidence. The banking connector's Federal Register function was specifically inspected and found limited to financial sanctions and export controls; it did not gain general public-health policy approval. All 35,123 non-GitHub focal observations and 362 additional hosted-source link contexts were closed.

The fixed repository lane returned 7,442 observations across 318 pages, with two new scope exclusions. Smithery returned 13,454 observations across 160 pages with no novel identity. Code search returned 8,403 observations across 217 actual pages: 8,402 exact prior observations and one changed RoB 2 instruction body. Full-body and difference review reconciled that revision to its existing component without widening approval. Final canonical comparison adds five units in five projects, producing an immutable baseline of 6,479 units with no removals or unresolved conflicts. The consecutive-zero counter remains zero; the stopping rule has not been met.

Expansion round 15

Round 15 is closed and positive. The 240 repository queries returned 7,442 observations across 318 pages; 80 Smithery queries returned 13,454 observations across 160 pages. These lanes found no novel identities. One Smithery timeout was retained and successfully retried. The 240 indexed queries yielded 254 focal pages: 244 HTTP 200 responses, five HTTP 403 responses and five HTTP 404 responses. Six initial timeouts, including two after HTTP 200 headers, were retained and fully recovered before parsing. Of 1,387 linked repository identities, 533 matched canonical records, 848 matched earlier directory screens and six received individual primary review.

The six initial projects produced 15 component dispositions: four Eligible, seven Hold and four Exclude. An authored link led to one additional Brazilian data collection. Its 72-module manifest was screened; 32 module families received selected implementation reviews and 40 remained outside the bounded research-component scope. Those selected reviews produced 46 component dispositions: 25 Eligible, 17 Hold, two Exclude, one functional duplicate and one discovery-only catalogue. These are component reviews within one project, not 72 approved projects. Renamed MJSP catalogue operations were reconciled to the existing SINESP subset. Specific limitations and holds retain migrated endpoints, incorrect period/geometry handling, uncertain numeric units, unsupported medicine/program claims, and stale embedded references. All 35,317 non-GitHub focal observations and meaningful authored source links were closed.

Code search completed 80 queries across 214 actual pages, returning 8,182 exact prior file observations and no novel bodies. Final canonical comparison adds 29 units across four projects, with no removals or unresolved conflicts, producing 6,508 units in the next immutable baseline. The consecutive-zero counter remains zero and the stopping rule remains unmet.

Expansion round 16

Round 16 is closed and positive. All fixed queries completed: 240 repository queries over 318 pages, 80 Smithery queries over 160 pages, 240 indexed-directory queries and 80 code queries. Repository search returned 7,443 observations and one new scope exclusion; Smithery returned 13,453 observations and three new listings, producing two narrow approvals and one exclusion. Indexed discovery retained 255 focal pages: 245 returned HTTP 200, five returned 403 and five returned 404. One failed initial fetch was retained and successfully retried. The 1,391 linked identities comprised 537 canonical matches, 845 previous directory dispositions, three new primary exclusions and six secondary purpose exclusions.

Fresh link review added a public-data catalogue and record-query adapter. The two hosted listing approvals cover a cited behavioral-science corpus and a registered arXiv metadata retrieval subset. An authored publisher follow-up produced 14 selected source/profile dispositions: six eligible components, one hold, six exclusions and one discovery catalogue. These cover named instrument/construct/citation retrieval and narrowly described measurement preparation, approximate calculations, descriptive correlations, preliminary annotation and source metadata. Documentation and illustrative outputs do not establish statistical correctness, scale validation, causal effects or clinical suitability. All 35,494 non-GitHub focal observations and meaningful authored links were closed.

The same publisher also exposed 39 tool declarations and a broader 312-card capability inventory. These received context/contract screening, separately from selected full-profile reviews. The broader inventory yielded 261 purpose exclusions, 38 specific missing-contract holds, ten planned-component observations and three partial referrals to exact reviewed tools, with no additional approvals. These are not 312 tested implementations or 38 public hold cards.

Code search retained 203 actual pages and 8,180 observations: 8,170 exact prior observations and ten newly reviewed bodies. One bounded evidence-synthesis instruction qualified, eight bodies reconciled to existing components and one was excluded. Ten linked guide records were reviewed separately: three qualified narrowly, four duplicated existing components and three were supporting context. Final canonical reconciliation adds 13 components across five new projects, with no removals or unresolved conflicts, producing 6,521 units in the next immutable baseline. The consecutive-zero counter remains zero; the stopping rule remains unmet.

Expansion round 17

Round 17 is the first fully closed zero-yield round. All 240 repository queries completed over 318 pages and all 80 Smithery queries completed over 160 pages, with no newly observed identities in those lanes. The 240 indexed queries produced 256 focal pages: 246 returned HTTP 200, five returned 403 and five returned 404, with no transport failures. Of 1,407 linked identity observations, 541 matched canonical records and 849 matched earlier directory dispositions. The 17 new observations comprised ten secondary purpose exclusions, three internal directory-path parser artifacts and four primary reviews: two excluded and two held for specific documented input/transport defects. Parser artifacts were not treated as actual GitHub repositories.

All 36,290 non-GitHub focal observations and meaningful source links were reconciled. Code search completed 80 queries across 217 actual pages, returning 8,411 observations: 8,409 exact prior observations and two newly reviewed bodies, one excluded and one duplicate. All 61 retained retry errors were resolved, with actual terminal-page boundaries verified. Canonical comparison found zero new units, zero new projects and zero removals; the next frozen baseline remains at 6,521 units. The consecutive-zero count is one. The stopping rule is not yet met: one further fully completed zero-yield round is required.

Expansion round 18

Round 18 is the second successive fully closed zero-yield round. It completed 240 repository queries over 318 pages, 80 Smithery queries over 160 pages, 240 indexed queries and 80 code queries. The repository lane returned only prior identities; three new Smithery listings were excluded. The indexed lane retained 255 focal pages: 245 HTTP 200, five HTTP 403 and five HTTP 404 responses, without transport failures. Among 1,404 linked identity observations, 541 matched canonical records and 852 matched prior directory dispositions. Eleven new observations received seven contextual dispositions, including one internal integration-path artifact, and four promoted source reviews with three exclusions and one hold. An additional linked sibling was excluded after full secondary README review with its declared primary unavailable; that sibling was already represented in the contextual set.

The 36,015 non-GitHub focal observations and meaningful source links were closed. A current Libgen README was compared with the retained original: scholarly provider scope was already documented, and the only substantive change restricted remote reading by default. It was reconciled as a duplicate, with the narrower deployment limitation retained. Code search produced 8,412 observations over 217 actual pages: 8,410 exact prior observations and two newly reviewed bodies, one excluded and one duplicate. All 26 retained retry errors were resolved. Request pacing and Retry-After handling changed without changing the fixed queries. Source-linked exporter and recursive citation-crawl components were held for specific source defects; their supporting helpers did not create new approvals. No candidate code was executed.

Final canonical comparison again found zero new units, zero new projects and zero removals, retaining 6,521 units. Rounds 17 and 18 therefore provide two successive completed zero-yield rounds under the revised source-specific plan. The search stopping criterion is met. The original unamended query matrix remains incomplete, the repository-partition cutoff applies only to that lane, and unavailable or unindexed sources remain outside the demonstrated coverage.

The final registry update reconciled against the post-round-18 baseline contains 6,227 active cards: 6,051 new cards and 176 retained active historical cards. Its workbook has 6,228 rows, including one historical held record. Thirteen historical cards have explicit updates. Final validation covers all 6,410 approved review references and preserves all 6,051 existing new-card IDs. All 25 workbook fields match the expected records, and the inspected final views are readable. The application production build and two server-rendered checks pass. These public-card counts are distinct from the 6,521 canonical project/component units.

A later audit found stale web-summary assertion references in two copied freeze helpers. The correct contemporary closure summaries had already been manually inspected and preserved by hash in the manifests. Direct checks reconfirmed the relevant query completion and zero pending work, and the next helper was corrected before use. No frozen data, canonical outcome or round result changed.

Interpretation and remaining work

Search observations, individual reviews, canonical projects, components and public cards have different denominators. Repeated search hits, identical instruction bodies, imported copies and hosted/source aliases must not be counted as separate discoveries by default. Conversely, approval of one component must not approve an entire collection or its held siblings. Automated metadata/homonym screens, exact-copy reconciliation and URL-role screening remain labeled separately from individual source reviews.

Eligibility records describe bounded documented capabilities and limitations. Adjacent scientific tools and necessary research infrastructure are distinguished from direct RWE workflows. Candidate programs were not installed or runtime-tested. Static source inspection can identify contradictions, but documentation-level eligibility is not evidence of analytical validity or clinical suitability. Historical approval can be withdrawn when stronger source evidence contradicts it; the original identity and review history remain retained.

Remaining coverage limitations include the original unamended query matrix and the need for future dated maintenance. The completed search process retains the original matrix limitation, source-access constraints and earlier code-log audit gap. Meeting the revised stopping rule describes this documented search process at a date; it does not guarantee an exhaustive inventory of a changing public ecosystem.